


Google Cloud Platform

Resource
Management



Resource Management

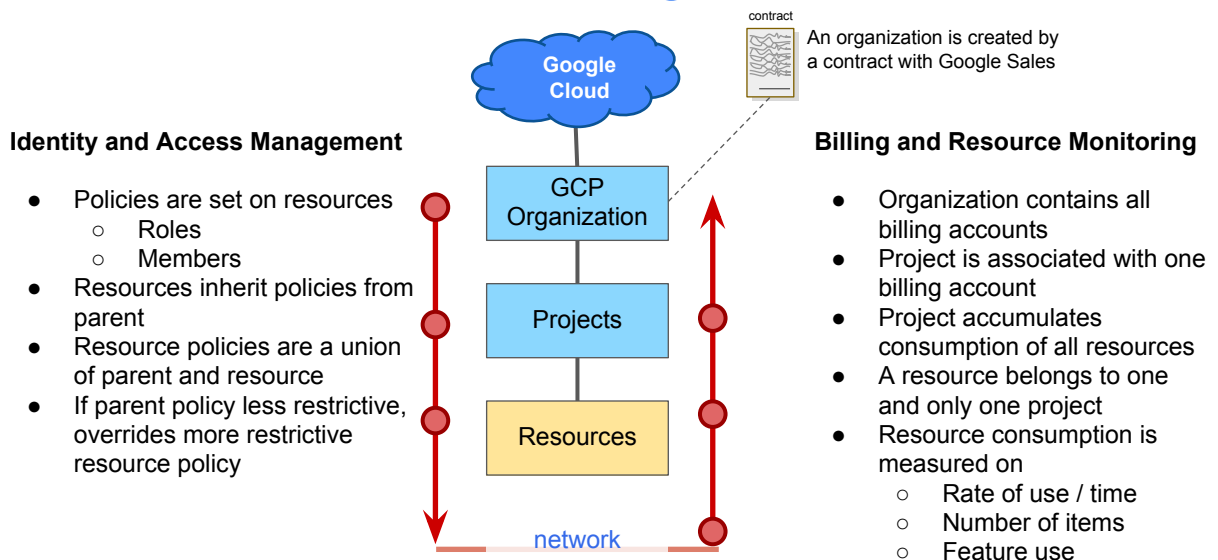
v 1.0



Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → Billing
- 5 → Lab
- 6 → Review


Cloud Resource Manager



- IAM Policies are inherited from the **top-down**
- IAM Policies determine which **users** (or groups) have which access **roles** to which **resources**
- Billing is accumulated from the **bottom-up**. Reporting of billing is per project.
- You can set policies at the organization, project, or individual resource level, to override policies set at the higher level in the hierarchy
- By default, all resources within the same Project are "trusted"
- There is isolation between Projects
 - Isolation is implemented in the network using Software Defined Networking (SDN), so that packets can be inspected and not permitted to pass between projects without authorization

Organization Node

- Organization node is root node for Google Cloud resources
- 2 organization roles:
 - *Organization Admin* - Control over all cloud resources
 - *Project Creator* - Controls project creation



Projects

- Track resource and quota usage
 - Enable billing
 - Manage permissions and credentials
 - Enable services and APIs
- Projects use three identifying attributes:
 - Project Name
 - Project Number
 - Project ID -- also known as Application ID
- Cloud Console or the Cloud Resource Manager API

API actions:

- Get a list of all projects associated with an account.
- Create new projects.
- Update existing projects.
- Delete projects.
- Undelete, or recover, projects that you don't want to delete.

Projects

<https://cloud.google.com/compute/docs/projects>

<https://cloud.google.com/compute/docs/networks-and-firewalls>

A network can belong to only one project. And an instance can attach to only one network.

A project can contain up to 5 networks and 24 CPUs. A single network can contain a maximum of 7000 instances.

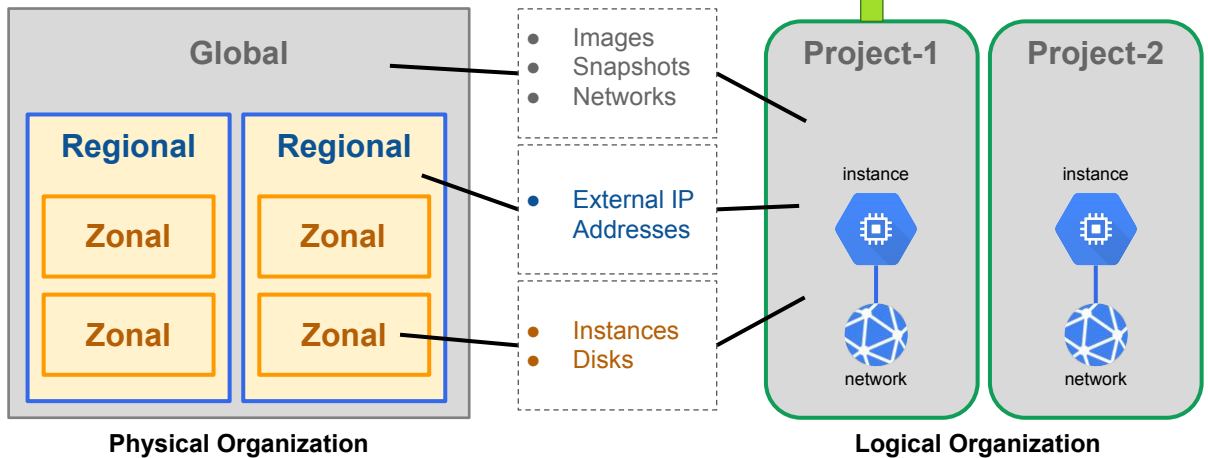
If you want servers to communicate over the Private IP, they have to be in the same project on the same network.

Since networks are homed to projects, as soon as two servers are located in different projects or on different networks in the same project, they must communicate over external IP addresses (which can be ephemeral or static) Internal DNS resolver of FQDN is part of Compute Engine.

<https://cloud.google.com/compute/docs/networking>

Resource hierarchy

Resources are Global, Regional or Zonal.



Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → Billing
- 5 → Lab
- 6 → Review

Project Quotas

- All resources are subject to project quotas or limits
 - Typically fall into one of three categories
 - How many resources you can create per project
 - How quickly you can make API requests in a project - rate limits
 - Some quotas are per region
 - Quota examples:
 - 5 networks per project
 - 24 CPUs region/project
 - Most quotas can be increased through self-service form or a support ticket

Compute Engine enforces quotas on resource usage for a variety of reasons. For example, quotas protect the community of Google Cloud Platform users by preventing unforeseen spikes in usage. Special quotas limit access for projects that are just exploring Google Cloud Platform on a free trial basis.

Not all projects have the same quotas. As your use of Google Cloud Platform expands over time, your quotas may increase accordingly. If you expect a notable upcoming increase in usage, you can proactively request quota adjustments from the Quotas page in the Cloud Platform Console.

How quotas are applied

Resource quotas are the maximum amount of resources you can create for that resource type, if those resources are available. Quotas do not guarantee that resources will be available at all times. If a resource is not available, you won't be able to create new resources of that type, even if you still have remaining quota in your region or project. This is particularly relevant for regional quotas; if a particular region is out of a resource, you won't be able to create a resource of that type, even if you still have quota. For example, if a region is out of local SSDs, you cannot create local SSDs in that region, even if you still had quota for local SSDs. In such cases, you should deploy regional resources in another region.

Check your quota

To check the available quota for resources in your project, go to the Quotas page in the Google Cloud Platform Console. If you are using gcloud, run the following command to check your quotas. Replace myproject with your own project ID:

```
gcloud compute project-info describe --project myproject
```

To check your used quota in a region, run:

```
gcloud compute regions describe example-region
```

Why Quotas?

- Prevent runaway consumption in case of an error or malicious attack
- Prevent billing spikes or surprises
- Enforce sizing consideration and periodic review

Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → Billing
- 5 → Lab
- 6 → Review

Labels

- A utility for organizing GCP resources
 - Attached to resources: VM, disk, snapshot, image
 - Console or API, currently not in `gcloud`

Example uses of labels:

- Search and list all resources (inventory)
- Filter resources (ex: separate production from test)


Labels used in scripts


- help analyze costs
- run bulk operations

<https://cloud.google.com/resource-manager/docs/using-labels>





Labels specification


- A label is a key:value pair
- Label keys and non-empty label values can contain lowercase letters, digits and hyphens, must start with a letter, and must end with a letter or digit. The regular expression is `[a-z]([-a-z0-9]*[a-z0-9])`
- The length of labels keys and values is up to 63 characters
- There can be a maximum of 64 labels per resource

 Labels

 Edit labels

Instances Selected (1): instance-1

Key	Value	
department	<input type="text" value="website-deve"/>	
engineering	<input type="text" value="development"/>	
owner	<input type="text" value="bobzalman"/>	
project	<input type="text" value="account-1569"/>	

 Add label

Labels can only be managed through the Compute Engine Beta API, or through the Google Cloud Platform Console. Labels are currently not supported through the gcloud command-line tool.

<https://cloud.google.com/compute/docs/label-or-tag-resources>

Label practices (1)

- Team or Cost Center
 - Distinguish projects owned by different teams.
 - Useful in cost accounting or budgeting.
 - Examples: `team:marketing`, `team:research`
- Components
 - Examples: `component:redis`, `component:frontend`
- Environment or stage
 - Examples: `environment:prod`, `environment:test`

Label practices (2)

- Owner or contact
- Person responsible for resource or primary contact for the resource
 - Examples: `owner:gaurav`, `contact:opm`
- State
 - Examples: `state:inuse`, `state:readyfordeletion`

Comparing Labels and Tags

- Labels are a way to organize resources across GCP
 - disks, image, snapshots...
- User-defined strings in `key:value` format
- Propagated through billing
- Tags are applied to instances only
- User-defined strings
- Tags are primarily used for networking (applying firewall rules)

Before March 7, 2017 VM tags were automatically created and synchronized with Labels, appearing as value-less Labels.

Create a tag, a label was created. Create a label, a tag was created. After that date, the two systems were separate -- two name spaces.

Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → **Billing**
- 5 → Lab
- 6 → Review

Budgets and Alerts

Billing Account or Selected Project

▼

Budget Amount

Specified Amount

▼

\$ 500.00

Last Month's Spend

Percent of Budget	Amount	
50	%	\$ 250.00 ✕
90	%	\$ 450.00 ✕
100	%	\$ 500.00 ✕

+ Add Item

Save

- Budget lets you track how spend is approaching specified amount
- Monthly budget (period is fixed)
- Can only be created by Billing Administrators
- Budget applies to either an entire billing account or to a single project
- Budget amount is specified or can be automatically set to last month's spend on the account or project
- Alerts are set as a % of budget (0.005% to 100%) rounds up to the cent
- Notification is triggered when spend is greater than alert amount (frequency ~hourly)
- Notification is sent by email to the Billing Administrator that created it (it does not appear in the console notifications)
- You can choose to include consumption of credits in the calculation or not (Credits are promotions or grants)

Example Notification Email

Billing Alert Notification

Dear Google customer,

You are receiving this email because you are a Google Cloud Platform, Firebase, or API customer.

This is an automated notification to inform you that the project: **deadpool-cpb100** has exceeded **0.05%** of the monthly budget of **\$100.00**.

You are receiving this message because there is an alert configured on this project's budget.

To disable this alert or modify the [budget's](#) threshold, please edit [your budget](#).

Billing Export

JSON Field	CSV Field	Data Type	Description
accountID	Account ID	string	Billing account ID
lineItemID	Line Item	string	URI of the resource
startTime	Start Time	dateTime	Start of measured period of use
endTime	End Time	dateTime	End of measured period of use
projectNumber	Project Number	integer	Project number
projectID	Project ID	string	Project ID
projectName	Project Name	string	Project Name
projectLabels	Project Labels	string	Project Labels
measurementID	Measurement	string	URI of the resource
sum	Measurement Total Consumption	integer	Measured time of use
unit	Measurement Units	string	Time period units (ie seconds)
creditID	Credit	string	Credit grant ID
amount	Credit Amount	decimal	Amount of the credit
currency	Credit Currency	string	Currency code (ie USD)
cost	Amount	decimal	Calculated cost
currency	Currency	string	Currency code (ie USD)

- Export to either a File or to a BigQuery dataset
- Create a GCS bucket or BQ dataset first, to specify when enabling
- Access is set via IAM on bucket or dataset
- File export is either in CSV or JSON format (not both)
- File export prefix name is appended with date-time-stamp
- Report is generated daily - there is no on demand generation
- Project Name and Project Labels are your primary post-export parsing tools

Fields that are exported to BigQuery:

<https://support.google.com/cloud/answer/7237695>

Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → Billing
- 5 → Lab
- 6 → Review

Lab #1

In this lab you will learn about billing administration inside GCP.

For security purposes, the lab environment may not be able to grant you the Billing Administrator role. The lab is designed to be a walk-through with visuals so that you can see the activities even if you can't perform them directly.

05-1 Billing Administration

Lab #2

BigQuery and Billing Data

In this lab you will create datasets and tables, import data from billing reports, and conduct a variety of queries on the data using BigQuery.

```
#standardSQL
```

```
SELECT
```

```
  billing_account_id
```

```
FROM
```

```
  `deadpool_billing_example.gcp_billing_export_006CDA_7A95F9_6F4862`
```


Agenda

- 1 → Cloud Resource Manager
- 2 → Quotas
- 3 → Labels and Names
- 4 → Billing
- 5 → Lab
- 6 → Review

More...

- Cloud Resource Manager
 - <https://cloud.google.com/resource-manager/>
- Quotas
 - <https://cloud.google.com/compute/quotas>
- Labels
 - <https://cloud.google.com/resource-manager/docs/using-labels>
 - <https://cloud.google.com/compute/docs/label-or-tag-resources>

More to learn on this subject. Here are some suggestions and links.

